

Exception: a transfer of control to the OS kernel in response to some event

Asynchronous Exceptions (Interrupts)

- Caused by events external to the processor

- Indicated by setting the processor's interrupt pin
- Handler returns to **"next" instruction**

Timer interrupt: take control back from user programs to kernel eg. context switch

Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:

- **Traps**

- Intentional
 - Examples: system calls (requests for services from the kernel)
 - Returns control to **"next" instruction**

- **Faults**

- Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable)
 - Either re-executes faulting (**"current"**) instruction or aborts

- **Aborts**

- Unintentional and unrecoverable
 - Examples: illegal instruction, parity error (data error/inconsistency detected), machine check (hardware issue detected)
 - **Aborts current program**

System Calls

Request a service that is not accessible for program from OS

- Each x86-64 system call has a unique ID number (assigned by the operating system)

- Examples:

Number	Name	Description
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

Process: an instance of a running program

Abstraction (Illusion):

■ Logical control flow


- Each program seems to have exclusive use of the CPU
- Provided by kernel mechanism called context switching

■ Private address space


- Each program seems to have exclusive use of main memory.
- Provided by kernel mechanism called virtual memory

Concurrent Processes

- Each process is a logical control flow.
- Two processes run **concurrently** (are concurrent) if their flows overlap in time
- Otherwise, they are **sequential**
- Examples (running on a single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C


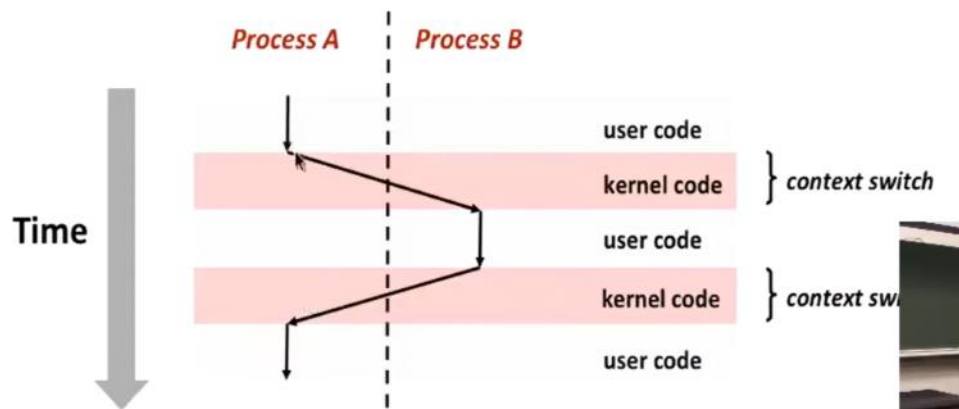


View of Concurrent Processes




Context Switching

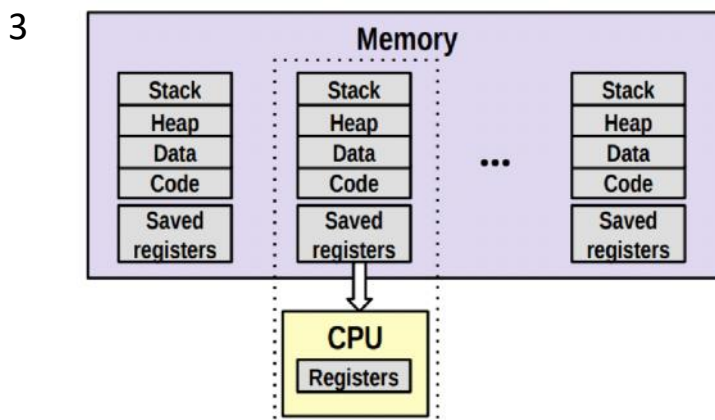
- Processes are managed by a shared chunk of memory-resident OS code called the **kernel**
 - Important: the kernel is not a separate process, but rather runs as part of some existing process.
- Control flow passes from one process to another via a **context switch**



- Save current registers in memory



- Schedule next process for execution



- Load saved registers and switch address space (**context switch**)

1. `void exit (int status)` - terminate the process, called once but never returns

■ Parent process creates a new running child process by calling `fork`

■ `int fork(void)`

- Returns 0 to the child process, child's PID to parent process
 - Child is **almost** identical to parent:
 - Child gets an identical (but separate) copy of the parent's virtual address space (this includes all the data on the stack and on the heap, and all the instructions)
 - Child gets identical copies of the parent's open file descriptors
 - Child has a different PID than the parent
- 2.

■ `fork(...)` function is interesting (and often confusing) because it is called once but returns twice

- 3.

Reaping Child Processes

■ Idea

- When process terminates, it still consumes system resources
 - Examples: Exit status, various OS tables
- Called a "**zombie**"
 - Living corpse, half alive and half dead

■ Reaping (harvesting, collecting)

- Performed by parent on terminated child process (using `wait` or `waitpid`)
- Parent is given exit status information (it is notified that the child process terminated and, by receiving the exit status, it acknowledges the termination)
- Kernel then deletes zombie child process

■ What if parent doesn't reap?

- If any parent terminates without reaping a child, then the orphaned child process will be reaped by init process (`pid == 1`) (*root of the process tree*)
- So, only need explicit reaping in long-running processes
 - e.g., shells and servers
- (although you should be a good citizen and collect your zombies if possible)

`wait`: Synchronizing with Children

■ Parent reaps a child by calling the `wait` function

■ `int wait(int *child_status)`

- suspends current process until one of its children terminates
- return value is the pid of the child process that terminated

only wait its own child (not grandchild)

Parent instructions after `wait` will be processed when child terminates

- 4.

execve : Loading and Running Programs

■ `int execve(char *filename, char *argv[], char *envp[])`

■ Loads and runs in the current process:

- Executable file filename
 - Can be object file or script file beginning with `#!/interpreter` (e.g., `#!/bin/bash`)
- ...with argument list `argv`
 - By convention `argv[0]==filename`
- ...and environment variable list `envp`
 - "name=value" strings (e.g., `USER=droh`)
 - `getenv`, `putenv`, `printenv`

■ Overwrites code, data, and stack

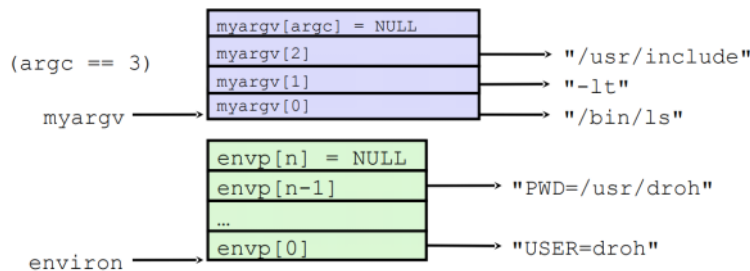
- Retains PID, open files and signal context
- (the current process is gone, it is now running a different program)

■ Called once and never returns

- ...except if there is an error

execve Example

■ Executes `"/bin/ls -lt /usr/include"` in child process using current environment:



```
if ((pid = Fork()) == 0) { /* Child runs program */  
    if (execve(myargv[0], myargv, environ) < 0) {  
        printf("%s: Command not found.\n", myargv[0]);  
        exit(1);  
    }  
}
```

5. `Kill(pid, SIGUSR1); kill(pid, SIGKILL); kill(pid, SIGSTOP)`
Send signal to pid